

Java PriorityQueue Class

Anders Jacobsen, Dima Karaush

May 6, 2021

Abstract

Java's PriorityQueue class have a bottleneck when you need to update values in the queue. Depending on the use and implementation, Java's PriorityQueue can add serious performance issues when accessing the data using the poll() method. When updating a value in the queue, performance can be improved by more than 50% compared to Java's PriorityQueue implementation. Implementing your own version might remove this bottleneck from your software.

1 Introduction

Why did we choose this subject This article will discuss Java's "Build-in" class `PriorityQueue`. The interest for this topic has grown from an implmentation of a weighted graph in Java, which we found had a possible shortcomming for our use-case. The shortcomming was that to update a value in the queue, we had to use a linear search (loop) through the queue to update a value within. That is fine on a small scale, but what if need to draw a graph of all the cities in the world? We believe that this has space for optimization and that is why this article explores this subject.

We are going to use a previously developed `Timer` class.

How will we work with the subject

2 Scope

What will be in this article What will not be included in this article

3 Problem

3.1 Background

3.2 Problem Statement

The questions we will try to solve Problemfomulering

4 Analysis

4.1 Benchmark Method

This subsection will cover how we prepared and executed our benchmarks. It will also touch the subject of what we are timing in the two implementations of a `PriorityQueue`. We've followed Peter Sestoft's approach to microbenchmarking in Java [1], we've used a combination of Mark5 and Mark3 benchmarks with a few twists here and there. To make benchmarks, a `Timer` class is necessary. We've designed the simplest version possible to avoid interference from calculations in the `Timer` class.

```
1 public class Timer {
2     private long start;
3
4     public void start() {
5         start = System.nanoTime();
6     }
7
8     public long step() {
9         return System.nanoTime() - start;
10    }
11 }
```

Listing 1: Simple `Timer` class implementation

This implementation enables us to process the times as nano seconds after the benchmarks and also to easily restart the timer.

In addition to the `Timer` class we have also implemented a `TimerTracker` class. This class only consists of two lists that can contain the warmup and real benchmark times. Also a method for writing the obtained times to a CSV¹ file. The CSV files will be used to explore the data later.

Measurements in the benchmark is done only on the time it takes to update a value in the `PriorityQueue`. Again we follow Peter Sestoft's microbenchmarking techniques. In listing 2 we run a number of warmup iterations before running the actual benchmark. This is to fight the battle against Java's JIT² compiler as described in Peter Sestoft's article. In the listing the method `pQueueRun()` is running the benchmark, this will be described in section 4.2. The `tracker.addTime(long time)` simply adds a time to a list that will later be written to a CSV file.

```
1 private static void benchmarkPriorityQueue(int
warmupIterations, int iterations, TimeTracker
tracker) {
2     // Printing removed for simplicity
3     // Warmup
```

¹Comma Separated Values Filestructure

²Just In Time

```

4         for (int i = 0; i < warmupIterations; i++) {
5             tracker.addWarmupTime(pQueueRun());
6         }
7
8         System.out.println("Warmup finished - moving
to real benchmark");
9
10        // Benchmark
11        for (int i = 0; i < iterations; i++) {
12            tracker.addTime(pQueueRun());
13        }
14    }

```

Listing 2: Benchmark iterations

4.2 Benchmark of Java's priorityQueue

To understand the example that we are using we have to take a look at how we use our queue.

```

1    PriorityQueue<Node> pQueue = new PriorityQueue<>(
    QUEUE_LENGTH, nodeComp);
2
3    // Filling the queue with Nodes
4    for (int i = 0; i < QUEUE_LENGTH; i++) {
5        pQueue.add(new Node(i, i + 1, i));
6    }

```

Listing 3: Populating the queue

We create a simple priorityQueue that takes nodes as an object. Then we add some nodes that fill up the queue, the nodes do not have any special values. But we do add a unique values so that we are able to distinguish the nodes.

When were looking at updating a node in Java's priority queue. There is no method to retrieve the wanted node from the list. Therefore we use a Iterator to list through every node in our list until we find a match.

```

1    Iterator<Node> it = pQueue.iterator();
2
3    while (it.hasNext()) {
4        n = it.next();
5        if (count == QUEUE_LENGTH - 1) {
6            n.verticeTo = 80085;
7            time = timer.step();
8            break;
9        }
10       count++;

```

11

}

Listing 4: Finding the node

Notice that we are listing through the queue in a linear way. This is the reason that we are updating the last object in the queue, to simulate a worst case scenario.

4.3 Update Method

4.4 Benchmark of updateable PriorityQueue

4.5 Comparisson of PriorityQueues

5 Conclusion

Answer our questions and possibly introduction

References

- [1] Peter Sestoft. Microbenchmarks in java and c-sharp. *IT University of Copenhagen, Denmark*, pages 2–16, 2015-09-16.